

INFORMED (HEURISTIC) SEARCH STRATEGIES

Course Code: **CSC4226** Course Title: **Artificial Intelligence and Expert System**



Dept. of Computer Science
Faculty of Science and Technology

Lecture No:	Four (4)	Week No:	Four (4)	Semester:	
Lecturer:	<i>Dr. Abdus Salam</i> <i>abdus.salam@aiub.edu</i>				

LECTURE OUTLINE



1. **BEST-FIRST SEARCH**
2. **GREEDY BEST-FIRST SEARCH**
3. **A* SEARCH**
4. **CONDITIONS FOR OPTIMALITY**
5. **OPTIMALITY OF A***

INFORMED SEARCH STRATEGY



Informed search strategy—

one that uses **problem-specific knowledge** beyond the **definition of the problem itself**

can find solutions more efficiently than can an uninformed strategy.

problem-specific knowledge is the extra bit of information the program uses rather than the problem formulation, thus known as **Informed Search**

BEST-FIRST SEARCH



- The general approach to informed search is called **best-first search**
- Best-first search is an instance of the general **TREE-SEARCH** Or **GRAPH-SEARCH** algorithm in which a node is selected for expansion based on an **evaluation function, $f(n)$** .
- The evaluation function is construed as a cost estimate, so the node with the *lowest* evaluation is expanded first
- The implementation of best-first graph search is identical to that for **uniform-cost search**, except for the **use of f instead of g** to order the priority queue.
- The choice of f determines the search strategy.

HEURISTIC SEARCH COMPARED WITH BLIND SEARCH



Brute force / Blind search

- ◆ Only have knowledge about already explored nodes
- ◆ No knowledge about how far a node is from goal state

Heuristic search

- ◆ Estimates "distance" to goal state
- ◆ Guides search process toward goal state
- ◆ Prefer states (nodes) that lead close to and not away from goal state

HEURISTIC FUNCTION

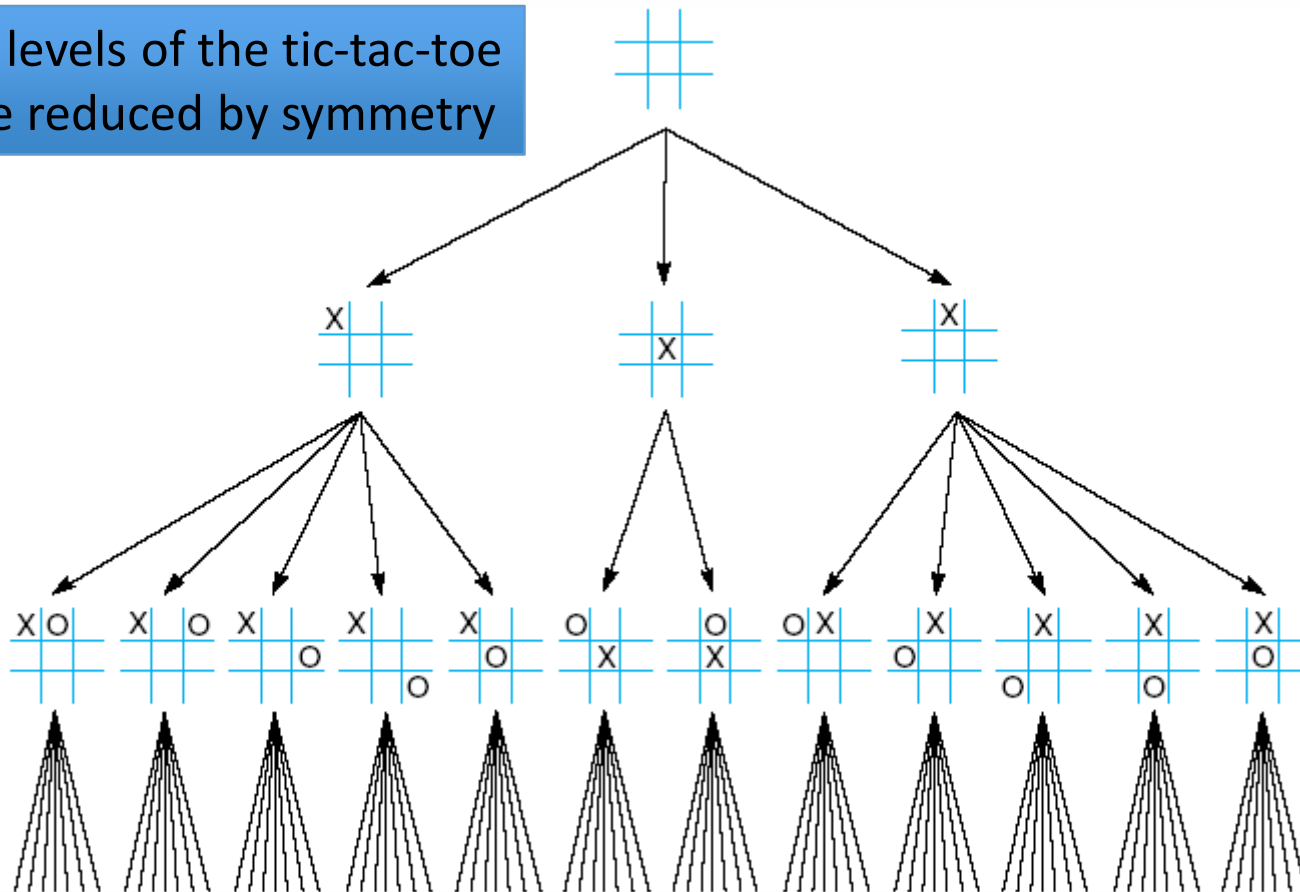


- The choice of **f(evaluation function)** determines the search strategy
- Best-first algorithms include as a component of **f** a **heuristic function**, denoted **$h(n)$**
- $h(n)$ = estimated cost of the cheapest path from the state at node n to a goal state.
- $h(n)$ takes a *node* as input, but, unlike $g(n)$, it depends only on the *state* at that node.
- Heuristic functions are the most common form in which additional knowledge of the problem is imparted to the search algorithm
- Consider them to be arbitrary, nonnegative, problem-specific functions, with one constraint: **if n is a goal node, then $h(n)=0$.**


HEURISTIC FUNCTION




First three levels of the tic-tac-toe state space reduced by symmetry




HEURISTIC FUNCTION



Three wins through
a corner square

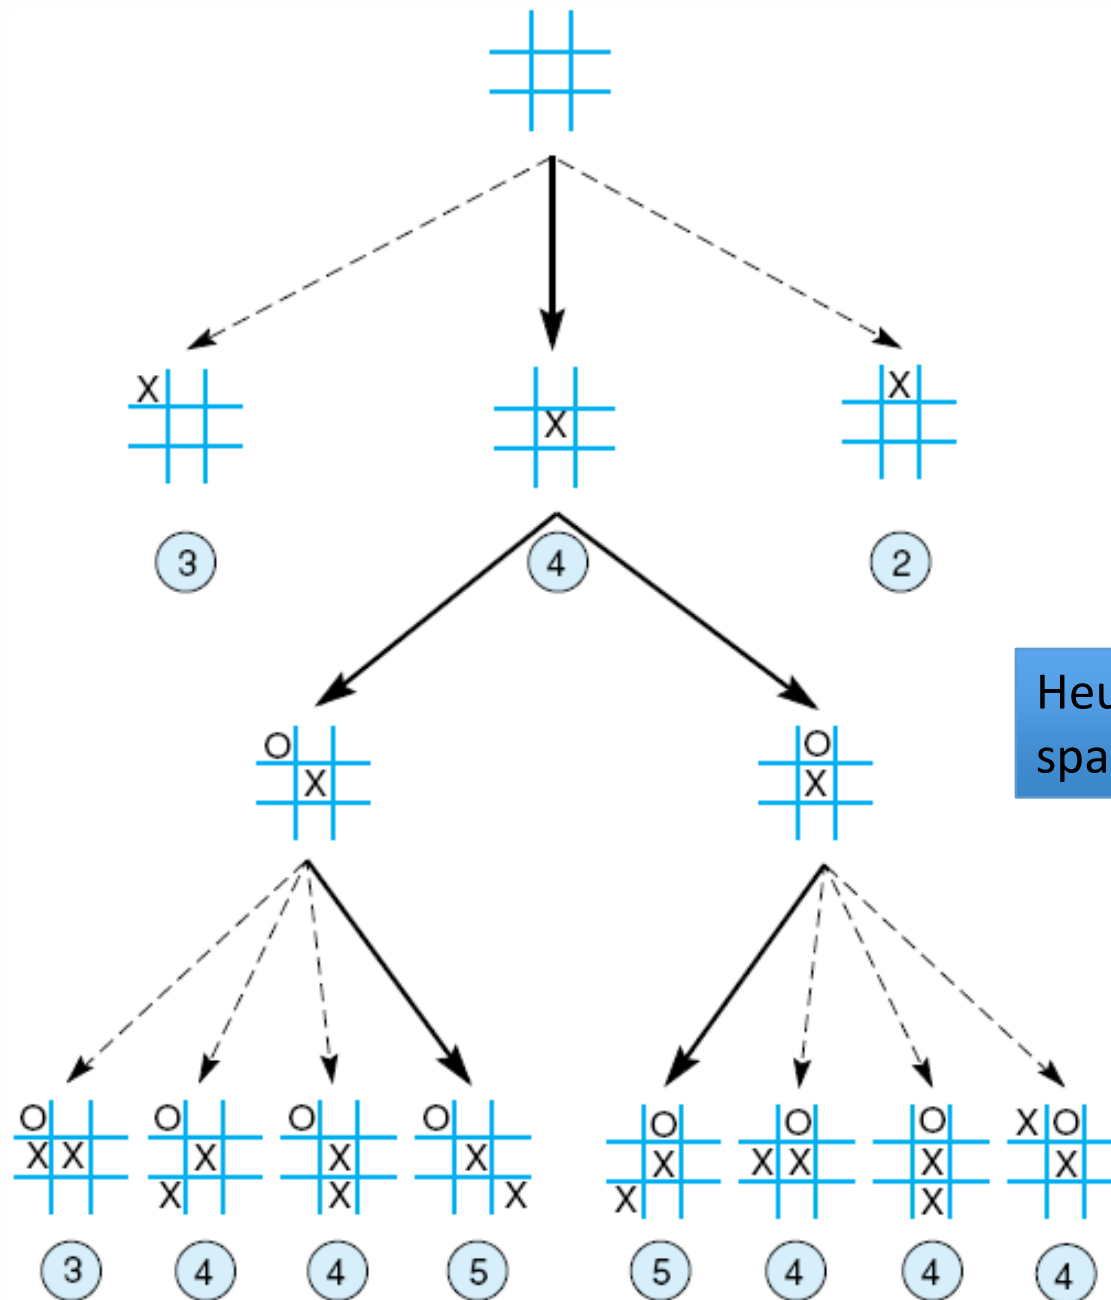


Four wins through
the center square



Two wins through
a side square

The “most wins” heuristic applied to the first children in tic-tac-toe.



Heuristically reduced state space for tic-tac-toe

GREEDY BEST-FIRST SEARCH



- **Greedy best-first search** tries to expand the node that is closest to the goal, on the grounds that this is likely to lead to a solution quickly.
- Thus, it evaluates nodes by using just the heuristic function; that is, $f(n) = h(n)$.
- “greedy”—at each step it tries to get as close to the goal as it can.
- Its search cost is minimal. It is **not optimal**.
- Greedy best-first tree search is also **incomplete** even in a finite state space, much like depth-first search.

STRAIGHT-LINE DISTANCE HEURISTIC



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Figure 3.16 Values of h_{SLD} —straight-line distances to Bucharest.


GREEDY BEST-FIRST TREE SEARCH




(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu



GREEDY BEST-FIRST TREE SEARCH



(d) After expanding Fagaras

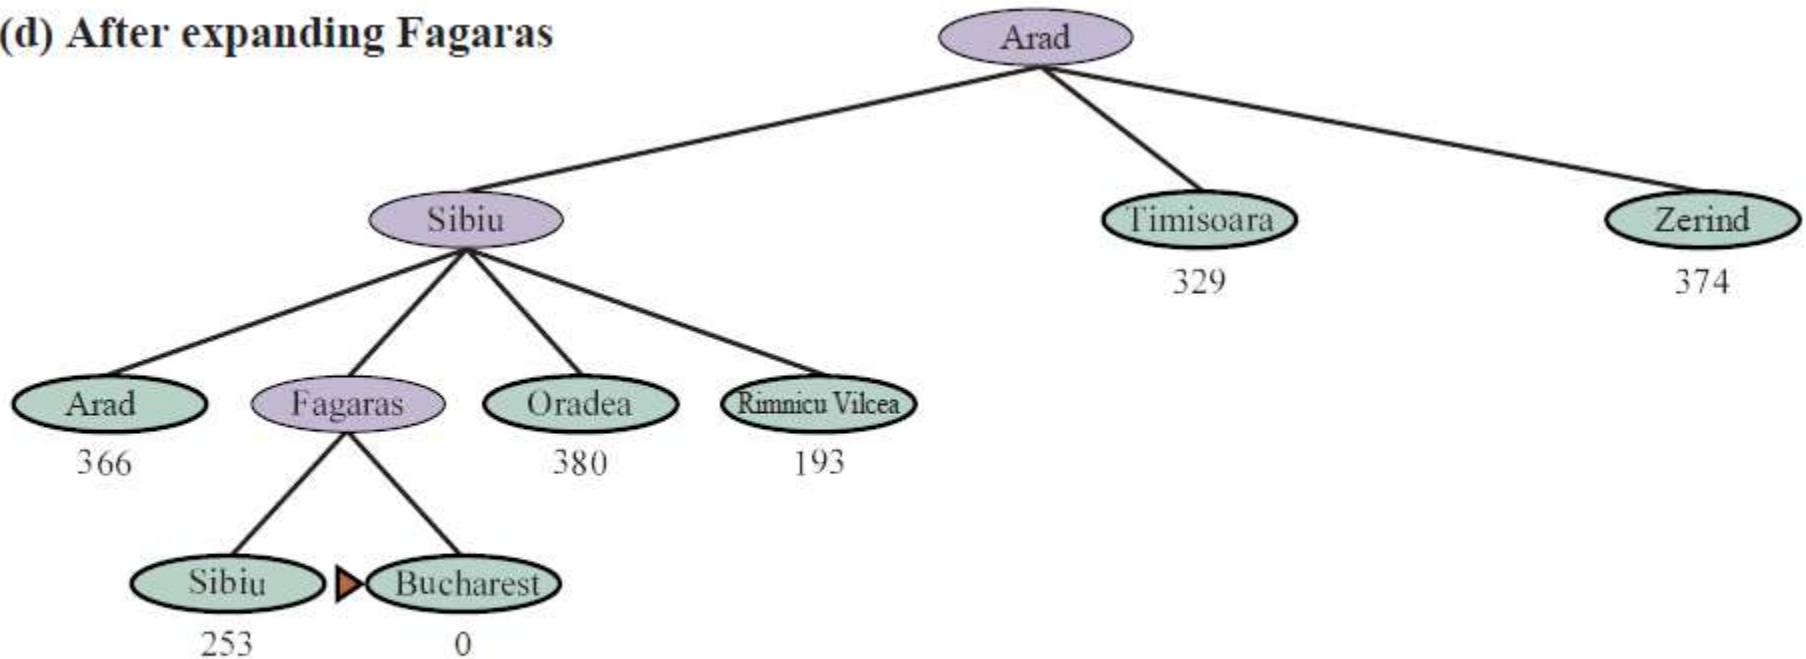


Figure 3.17 Stages in a greedy best-first tree-like search for Bucharest with the straight-line distance heuristic h_{SLD} . Nodes are labeled with their h -values.

A* SEARCH:

MINIMIZING THE TOTAL ESTIMATED SOLUTION COST



The most widely known form of best-first search is called **A* search**

It evaluates nodes by combining

$g(n)$, the cost to reach the node, and

$h(n)$, the cost to get from the node to the goal:

$$f(n) = g(n) + h(n) .$$

Since

$g(n)$ gives the path cost from the start node to node n , and

$h(n)$ is the estimated cost of the cheapest path from n to the goal,

we have $f(n)$ = estimated cost of the cheapest solution through n .

A* search is both **complete and optimal**.

The algorithm is identical to UNIFORM-COST-SEARCH except that

A* uses **$g + h$** instead of **g** .


A* SEARCH FOR BUCHAREST




(a) The initial state



(b) After expanding Arad



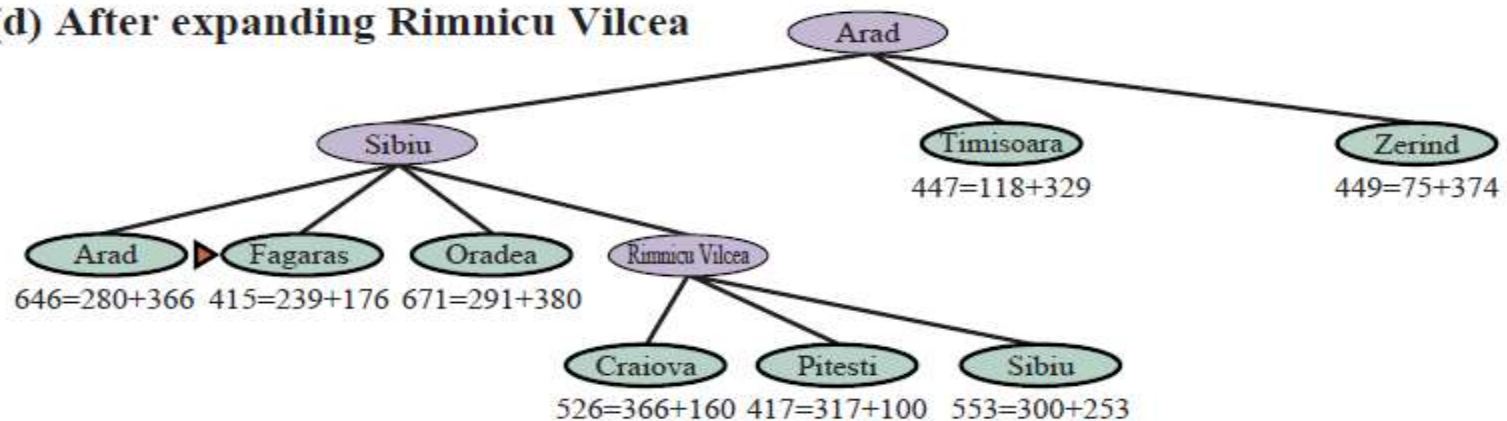
(c) After expanding Sibiu



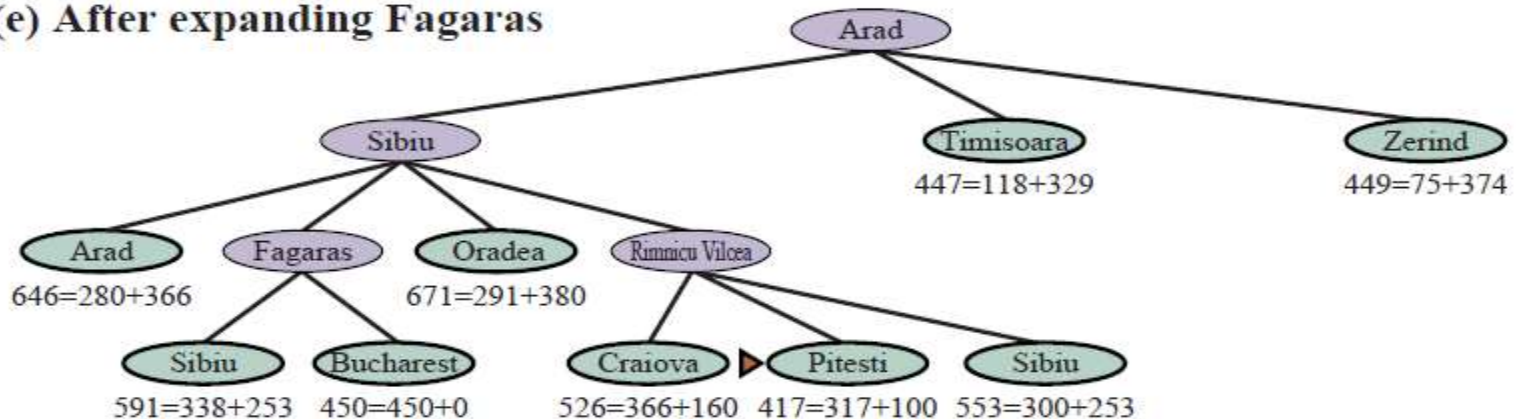
A* SEARCH FOR BUCHAREST



(d) After expanding Rimnicu Vilcea



(e) After expanding Fagaras



A* SEARCH FOR BUCHAREST



(f) After expanding Pitesti

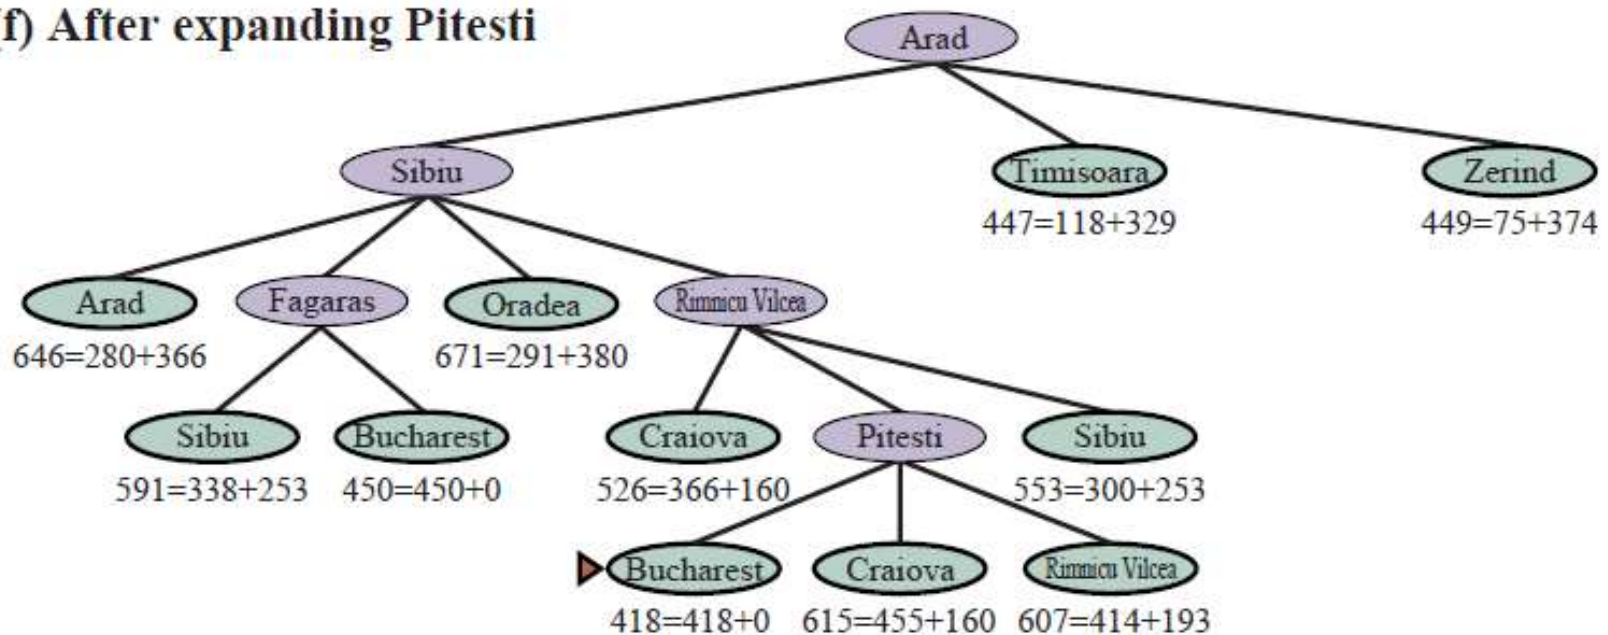


Figure 3.18 Stages in an A* search for Bucharest. Nodes are labeled with $f = g + h$. The h values are the straight-line distances to Bucharest taken from Figure ??.

CONDITIONS FOR OPTIMALITY: ADMISSIBILITY AND CONSISTENCY




- The first condition we require for optimality is that $h(n)$ be an **admissible heuristic**
- An admissible heuristic is one that *never overestimates* the cost to reach the goal
- Because $g(n)$ is the actual cost to reach n along the current path, and $f(n)=g(n) + h(n)$, we have as an immediate consequence that $f(n)$ never overestimates the true cost of a solution along the current path through n .
- Admissible heuristics are by nature optimistic because they think the cost of solving the problem is less than it actually is.

ADMISSIBILITY



- What must be true about h for A^* to find optimal path?
- A^* finds optimal path if h is admissible; h is **admissible** when it never overestimates.
- In this example, h is not admissible.
- In route finding problems, straight-line distance to goal is admissible heuristic.



$$g(X) + h(X) = 102$$

$$g(Y) + h(Y) = 74$$

Optimal path is not found!

CONSISTENCY



- A second, slightly stronger condition called **consistency** (or sometimes **monotonicity**) is required only for applications of A^* to graph search.
- A heuristic $h(n)$ is consistent if, for every node n and every successor n' of n generated by any action a , the estimated cost of reaching the goal from n is no greater than the step cost of getting to n' plus the estimated cost of reaching the goal from n' :

$$h(n) \leq c(n, a, n') + h(n').$$

- This is a form of the general **triangle inequality**, which stipulates that each side of a triangle cannot be longer than the sum of the other two sides

OPTIMALITY OF A*



A* has the following properties: *the tree-search version of A* is optimal if $h(n)$ is admissible, while the graph-search version is optimal if $h(n)$ is consistent.*

if $h(n)$ is consistent, then the values of $f(n)$ along any path are nondecreasing.

The proof follows directly from the definition of consistency. Suppose n' is a successor of n ; then $g(n') = g(n) + c(n, a, n')$ for some action a , and we have

$$f(n') = g(n') + h(n')$$

$$= g(n) + c(n, a, n') + h(n') \text{ because } g(n') = g(n) + c(n, a, n')$$

$$\geq g(n) + h(n) \text{ because } c(n, a, n') + h(n') \geq h(n)$$

$$= f(n)$$

$$\text{So, } f(n') \geq f(n)$$

OPTIMALITY OF A*



The next step is to prove that *whenever A^* selects a node n for expansion, the optimal path to that node has been found*. Were this not the case, there would have to be another frontier node n' on the optimal path from the start node to n ; because f is nondecreasing along any path, n' would have lower f -cost than n and would have been selected first.

Local search algorithms



- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- keep a single "current" state, try to improve it

Hill-climbing search



- "Like climbing Everest in thick fog with amnesia"

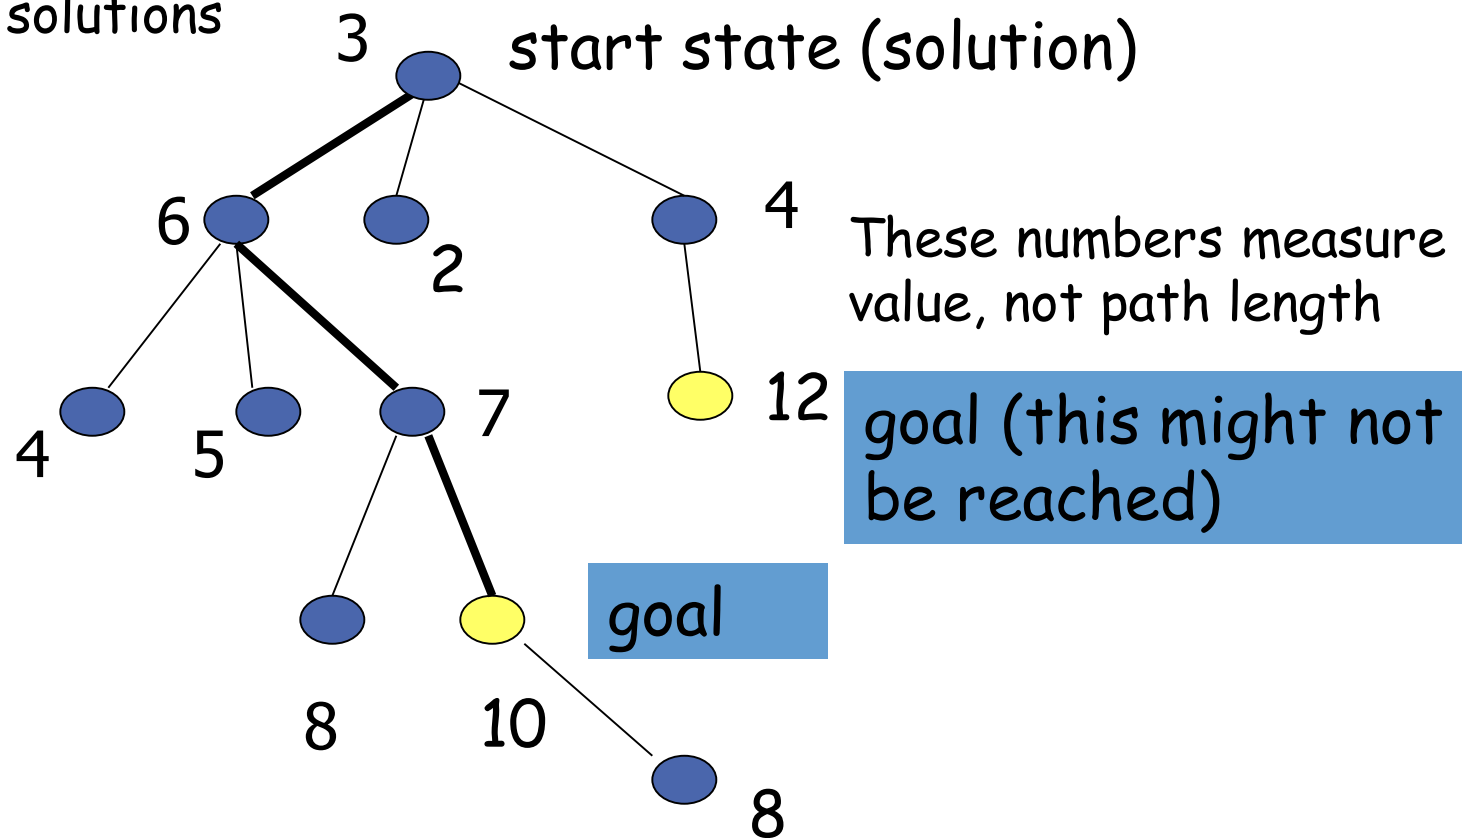
```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```


Example of hill-climbing




Neighbor solutions



Hill-climbing search




- Problem: depending on initial state, can get stuck in local maxima



Hill-climbing search



► Local maximum problem




Hill-climbing search



- ▶ **Ridge:** Result in a sequence of local maxima that is very difficult for a greedy algorithms to navigate.
- ▶ **Ridge:** then jump a little bit.


Ridge



Hill-climbing search



- ▶ **Plateau:** The evaluation function is flat. It can be a flat local maximum, from which no uphill exit exists or a shoulder, from which it is possible to make progress.
- ▶ **Plateau:** may be side way move. However, may not be a solⁿ when no uphill exists. So need some limitation of steeping sideway.





References

1. Chapter 3: Solving Problem by Searching , Pages 92-97
“Artificial Intelligence: A Modern Approach,” by Stuart J. Russell and Peter Norvig,



Books

1. "Artificial Intelligence: A Modern Approach," by Stuart J. Russell and Peter Norvig.
2. "Artificial Intelligence: Structures and Strategies for Complex Problem Solving", by George F. Luger, (2002)
3. "Artificial Intelligence: Theory and Practice", by Thomas Dean.
4. "AI: A New Synthesis", by Nils J. Nilsson.
5. "Programming for machine learning," by J. Ross Quinlan,
6. "Neural Computing Theory and Practice," by Philip D. Wasserman, .
7. "Neural Network Design," by Martin T. Hagan, Howard B. Demuth, Mark H. Beale, .
8. "Practical Genetic Algorithms," by Randy L. Haupt and Sue Ellen Haupt.
9. "Genetic Algorithms in Search, optimization and Machine learning," by David E. Goldberg.
10. "Computational Intelligence: A Logical Approach", by David Poole, Alan Mackworth, and Randy Goebel.
11. "Introduction to Turbo Prolog", by Carl Townsend.